

main

August 25, 2024

```
[ ]: import pandas as pd
import numpy as np
import plotly.express as px
import plotly.graph_objects as go
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression

from sklearn.metrics import mean_absolute_percentage_error
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score

import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
```

Data Description: Data Source: - This data is acquired from <https://docs.google.com/spreadsheets/d/1caaR9pT24GNmq3rDQpMiIMJrmiTGarbs/edit?gid=1150341366#gid=1150341366> as .xlsx file and converted to data.csv file. - This dataset contains the following columns: 1. Id: To count the records. 2. MSSubClass: Identifies the type of dwelling involved in the sale. 3. MSZoning: Identifies the general zoning classification of the sale. 4. LotArea: Lot size in square feet. 5. LotConfig: Configuration of the lot 6. BldgType: Type of dwelling 7. OverallCond: Rates the overall condition of the house 8. YearBuilt: Original construction year 9. YearRemodAdd: Remodel date (same as construction date if no remodeling or additions). 11. BsmtFinSF2: Type 2 finished square feet. 12. TotalBsmtSF: Total square feet of basement area 13. SalePrice: To be predicted Data Question: - What factors are significant in predicting sale price of houses? And what price is the house based on given factors's information? The response variable: SalePrice The possible predictors: MSSubClass, MSZoning, LotArea, LotConfig, BldgType, OverallCond, YearBuilt, YearRemodAdd, BsmtFinSF2, TotalBsmtSF

```
[ ]: # Load and show the first 10 rows of the data to see data's structure
df = pd.read_csv('data.csv')
df.head(10)
```

```
[ ]:   Id  MSSubClass MSZoning  LotArea LotConfig BldgType  OverallCond  \
0    0           60      RL     8450     Inside     1Fam           5
```

1	1	20	RL	9600	FR2	1Fam	8
2	2	60	RL	11250	Inside	1Fam	5
3	3	70	RL	9550	Corner	1Fam	5
4	4	60	RL	14260	FR2	1Fam	5
5	5	50	RL	14115	Inside	1Fam	5
6	6	20	RL	10084	Inside	1Fam	5
7	7	60	RL	10382	Corner	1Fam	6
8	8	50	RM	6120	Inside	1Fam	5
9	9	190	RL	7420	Corner	2fmCon	6

	YearBuilt	YearRemodAdd	Exterior1st	BsmtFinSF2	TotalBsmtSF	SalePrice
0	2003	2003	VinylSd	0.0	856.0	208500.0
1	1976	1976	MetalSd	0.0	1262.0	181500.0
2	2001	2002	VinylSd	0.0	920.0	223500.0
3	1915	1970	Wd Sdng	0.0	756.0	140000.0
4	2000	2000	VinylSd	0.0	1145.0	250000.0
5	1993	1995	VinylSd	0.0	796.0	143000.0
6	2004	2005	VinylSd	0.0	1686.0	307000.0
7	1973	1973	HdBoard	32.0	1107.0	200000.0
8	1931	1950	BrkFace	0.0	952.0	129900.0
9	1939	1950	MetalSd	0.0	991.0	118000.0

```
[ ]: # Check data type of each variable an correct the data type
df.info()
# Count categorical columns
categorical_columns = df.select_dtypes(include=['object', 'category', 'bool'])
print(f'Number of categorical columns: {len(categorical_columns.columns)}')
print(categorical_columns.columns)

# Count numeric columns:
numeric_columns = df.select_dtypes(include=['number'])
print(f'Number of numeric columns: {len(numeric_columns.columns)}')
print(numeric_columns.columns)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2919 entries, 0 to 2918
Data columns (total 13 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Id               2919 non-null   int64
1   MSSubClass       2919 non-null   int64
2   MSZoning         2915 non-null   object
3   LotArea          2919 non-null   int64
4   LotConfig        2919 non-null   object
5   BldgType         2919 non-null   object
6   OverallCond      2919 non-null   int64
7   YearBuilt        2919 non-null   int64
8   YearRemodAdd     2919 non-null   int64
```

```

9 Exterior1st    2918 non-null    object
10 BsmtFinSF2    2918 non-null    float64
11 TotalBsmtSF   2918 non-null    float64
12 SalePrice     1460 non-null    float64
dtypes: float64(3), int64(6), object(4)
memory usage: 296.6+ KB
Number of categorical columns: 4
Index(['MSZoning', 'LotConfig', 'BldgType', 'Exterior1st'], dtype='object')
Number of numeric columns: 9
Index(['Id', 'MSSubClass', 'LotArea', 'OverallCond', 'YearBuilt',
       'YearRemodAdd', 'BsmtFinSF2', 'TotalBsmtSF', 'SalePrice'],
      dtype='object')

```

- This dataset has 2919 entries and 13 columns.
- That includes: 4 categorical variables: 'MSZoning', 'LotConfig', 'BldgType', 'Exterior1st'
9 numeric variables: 'Id', 'MSSubClass', 'LotArea', 'OverallCond', 'YearBuilt', 'YearRemodAdd', 'BsmtFinSF2', 'TotalBsmtSF', 'SalePrice'

```

[ ]: # Drop Id column
if 'Id' in df.columns:
    df.drop('Id', axis=1, inplace=True)

# Check data statistic of numeric variables
df.describe()

```

```

[ ]:
      MSSubClass      LotArea  OverallCond  YearBuilt  YearRemodAdd  \
count  2919.000000    2919.000000    2919.000000    2919.000000    2919.000000
mean    57.137718    10168.114080      5.564577    1971.312778    1984.264474
std     42.517628     7886.996359      1.113131     30.291442     20.894344
min     20.000000     1300.000000      1.000000    1872.000000    1950.000000
25%     20.000000     7478.000000      5.000000    1953.500000    1965.000000
50%     50.000000     9453.000000      5.000000    1973.000000    1993.000000
75%     70.000000    11570.000000      6.000000    2001.000000    2004.000000
max     190.000000   215245.000000      9.000000    2010.000000    2010.000000

      BsmtFinSF2  TotalBsmtSF      SalePrice
count  2918.000000    2918.000000    1460.000000
mean    49.582248    1051.777587    180921.195890
std    169.205611     440.766258     79442.502883
min      0.000000      0.000000     34900.000000
25%      0.000000     793.000000    129975.000000
50%      0.000000     989.500000    163000.000000
75%      0.000000    1302.000000    214000.000000
max    1526.000000    6110.000000    755000.000000

```

```

[ ]: # Check missing/null values of data
df.isnull().sum()

```

```
[ ]: MSSubClass      0
      MSZoning       4
      LotArea        0
      LotConfig      0
      BldgType       0
      OverallCond    0
      YearBuilt      0
      YearRemodAdd   0
      Exterior1st    1
      BsmtFinSF2     1
      TotalBsmtSF    1
      SalePrice     1459
      dtype: int64
```

- MSZoning has 4 null values
- Exterior1st has 1 null values
- BsmtFinSF2 has 1 null values
- TotalBsmtSF has 1 null values
- SalePrice has 1459 null values

```
[ ]: # Replace null values with mean values and remove duplicate rows if any in the
      ↪dataset
df. drop_duplicates(inplace=True)

# Fill na for categorical variables
df['MSZoning'].fillna(df['MSZoning'].mode()[0], inplace = True)
df['Exterior1st'].fillna(df['Exterior1st'].mode()[0], inplace = True)

# Fill na for numeric variables
df['BsmtFinSF2'].fillna(df['BsmtFinSF2'].mean(), inplace = True)
df['TotalBsmtSF'].fillna(df['TotalBsmtSF'].mean(), inplace = True)
df['SalePrice'].fillna(df['SalePrice'].mean(), inplace = True)
df.shape
```

```
[ ]: (2911, 12)
```

After cleaning: - This dataset has 2911 entries and 12 columns. - That includes: 4 categorical variables: 'MSZoning', 'LotConfig', 'BldgType', 'Exterior1st' 9 numeric variables: 'MSSubClass', 'LotArea', 'OverallCond', 'YearBuilt', 'YearRemodAdd', 'BsmtFinSF2', 'TotalBsmtSF', 'SalePrice'

```
[ ]: # Plot bar chart to see the distribution of each variable

# Plot categorical variables:
for col in categorical_columns:
    unique_values = df[col].unique()
    print(f'Values of {col}: ', unique_values)
    value_count = df[col].value_counts().reset_index()
    value_count.columns = ['Category', 'Count']
```

```

    plot = px.bar(value_count, x='Category', y='Count', title=f'{col} Category_
↪Distribution')
    plot.show()

```

Values of MSZoning: ['RL' 'RM' 'C (all)' 'FV' 'RH']

Values of LotConfig: ['Inside' 'Corner' 'FR2' 'CulDSac' 'FR3']

Values of BldgType: ['1Fam' 'TwnhsE' 'Duplex' 'Twnhs' '2fmCon']

Values of Exterior1st: ['VinylSd' 'Wd Sdng' 'HdBoard' 'BrkFace' 'WdShing'
'MetalSd' 'CemntBd'
'Stucco' 'Plywood' 'AsbShng' 'Stone' 'ImStucc' 'CBlock' 'BrkComm']

```

[ ]: # Plot Numerical Variables
numeric_columns = df.select_dtypes(include=['number'])
for col in numeric_columns:
    fig = px.histogram(df, x=col, nbins=30, title=f'Histogram of {col}')
    # Customize the appearance if needed
    fig.update_layout(
        xaxis_title=col,
        yaxis_title='Counts',
        bargap=0.1, # Adjust the gap between bars
        title_x=0.5 # Center the title
    )
    fig.show()

```

```

[ ]: # Explore relationship between numerical variables


# Select only numerical columns
numerical_df = df.select_dtypes(include='number')

# Compute the correlation matrix for the numerical columns
correlation_matrix = numerical_df.corr()

# Plotting the heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(correlation_matrix, annot=True, fmt='.2f', cmap='coolwarm',
↪cbar=True)

# Update layout for better readability
plt.title('Correlation Matrix Heatmap')
plt.show()

```



- TotalBsmtSF and SalePrice have a positive correlation (~ 0.43), indicating that as the total basement area increases, the sale price tends to increase.
- YearBuilt, YearRemodAdd also show a positive correlation ($0.35\sim 0.37$) with SalePrice, suggesting that newer houses tend to have higher sale prices.
- OverallCond, MSSubclass have a slight negative correlation with SalePrice ($-0.05\sim -0.06$), meaning that the overall condition rating may slightly decrease as sale price increases, though the effect is very minimal.
- LotArea and SalePrice have a correlation of 0.24 suggests that, in general, larger lot areas tend to be associated with higher sale prices. However, the relationship is not very strong, implying that other factors likely play a more significant role in determining the sale price of a property.

```
[ ]: # Explore relationship between SalePrice and other categorical variables using correlation ratio

# Function to calculate the correlation ratio (²)
def correlation_ratio(categories, measurements):
    fcat, _ = pd.factorize(categories)
```

```

cat_num = np.max(fcat) + 1
y_avg_array = np.zeros(cat_num)
n_array = np.zeros(cat_num)

for i in range(cat_num):
    cat_measures = measurements[fcat == i]
    n_array[i] = len(cat_measures)
    y_avg_array[i] = np.mean(cat_measures)

y_total_avg = np.sum(y_avg_array * n_array) / np.sum(n_array)
numerator = np.sum(n_array * (y_avg_array - y_total_avg) ** 2)
denominator = np.sum((measurements - y_total_avg) ** 2)

if numerator == 0:
    return 0.0
else:
    return np.sqrt(numerator / denominator)

# Calculate the correlation ratio for each categorical variable
for col in categorical_columns:
    # Ensure there are no missing values in the column and SalePrice
    eta_squared = correlation_ratio(df[col].dropna(), df['SalePrice'].
    ↪loc[df[col].dropna().index])
    print(f"Correlation Ratio ( 2 ) for {col}: {eta_squared}")

```

```

Correlation Ratio ( 2 ) for MSZoning: 0.22669722318599306
Correlation Ratio ( 2 ) for LotConfig: 0.10523387811917098
Correlation Ratio ( 2 ) for BldgType: 0.12899402011999697
Correlation Ratio ( 2 ) for Exterior1st: 0.27626786820716165


```

- ² between 0.1 and 0.3: There is a weak to moderate association between categorical variables (MSZoning, LotConfig, BldgType, Exterior1st) and numeric variables (SalePrice). The numeric values show some variation across categories, but the relationship is not strong.

```

[ ]: # Detect outliers using Boxplot
feature = ['YearBuilt', 'YearRemodAdd', 'TotalBsmtSF', 'OverallCond',
    ↪ 'MSSubClass']
plt.figure(figsize=(15,15))
for i,col in enumerate(feature):
    plt.subplot(3,2,i+1)
    plt.boxplot(df[col])
    plt.title(f'{col}')
plt.show()

```



- The graph shows there are too some outliers in MSSubClass, OverallCond.

```
[ ]: # Handle outliers using Z-score

# Calculate Z-scores
z_scores = np.abs(stats.zscore(df[['YearBuilt', 'TotalBsmtSF', 'OverallCond', 'MSSubClass'])))

# Set a threshold for Z-score (common choice is 3)
threshold = 2


# Identify outliers
outliers = np.where(z_scores > threshold)
```



```
# Remove outliers
df = df[(z_scores < threshold).all(axis=1)]

# Reset index
df = df.reset_index(drop=True)
```

```
[ ]: # Boxplot after handling outliers
feature = ['YearBuilt', 'TotalBsmtSF', 'OverallCond', 'MSSubClass']
plt.figure(figsize=(15,10))
for i,col in enumerate(feature):
    plt.subplot(2,2,i+1)
    plt.boxplot(df[col])
    plt.title(f'{col}')
plt.show()
```



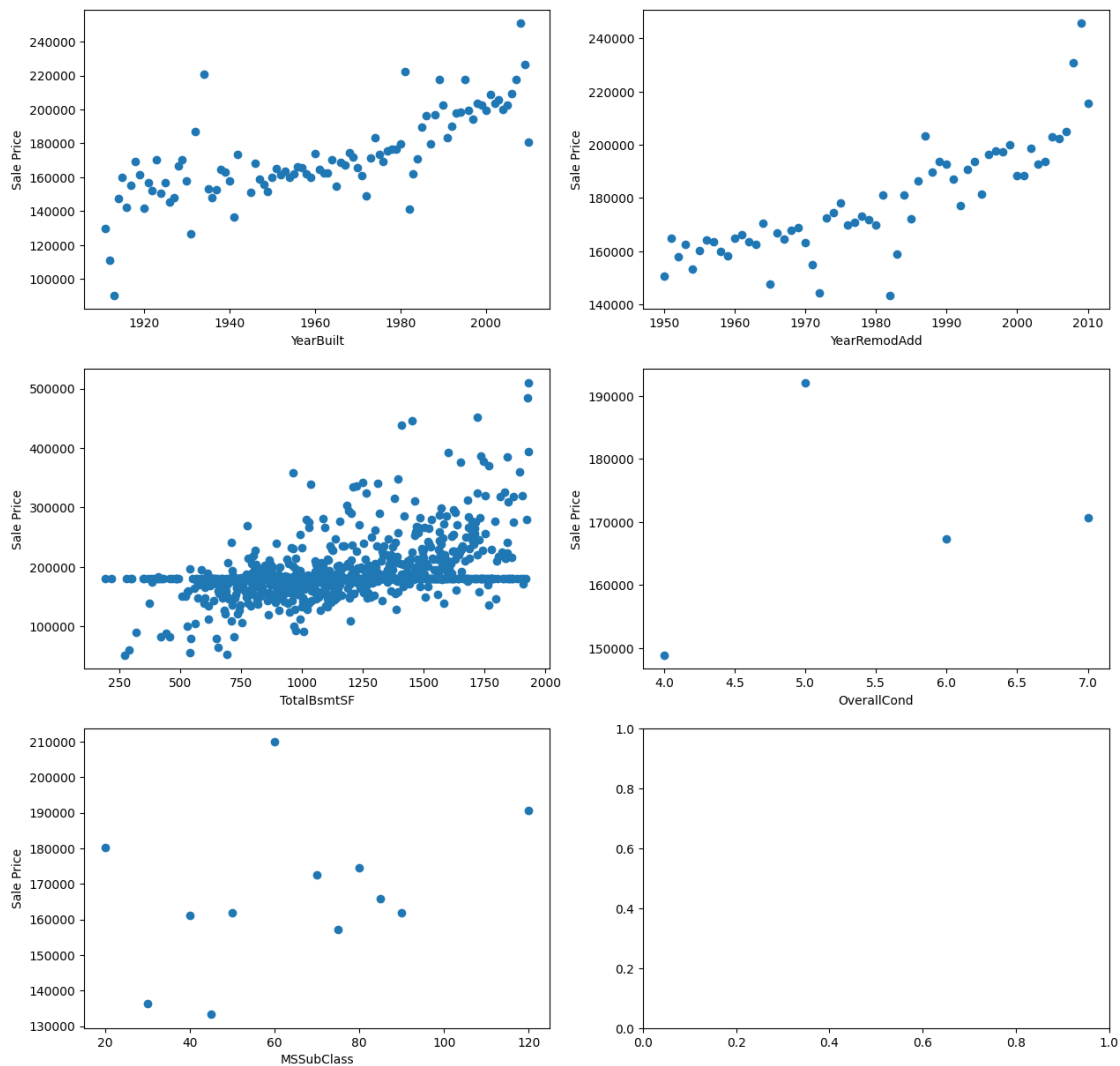
All of the outliers are removed from the dataset.

```
[ ]: # Plot the graphs to see relationships between SalePrice and 5 significant
      ↪ variables
feature = ['YearBuilt', 'YearRemodAdd', 'TotalBsmtSF', 'OverallCond',
      ↪ 'MSSubClass']
plt.subplots(3,2, figsize=(15,15))
```

```

for i, col in enumerate(feature):
    grouped_data = df.groupby(col)['SalePrice'].mean()
    plt.subplot(3,2,i+1)
    plt.scatter(grouped_data.index, grouped_data.values)
    plt.xlabel(f'{col}')
    plt.ylabel('Sale Price')
plt.show()

```



Based on the graphs, I decided to use linear models to build a ML model to predict the house's prices.

```

[ ]: # Build Linear Regression Model
X = df.drop(['SalePrice', 'MSZoning', 'LotArea', 'LotConfig', '
↳ 'BldgType', 'Exterior1st', 'BsmtFinSF2'], axis= 1)

```

```

Y = df['SalePrice']

# Split the data into training and testing sets
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.3,
    random_state=50)

# Fit a linear regression model
model = LinearRegression()
model.fit(X_train, Y_train)

# Predict on the test set
Y_pred = model.predict(X_test)

```

```

[ ]: # Calculate residuals
residuals = Y_test - Y_pred
# Evaluate model

mae = mean_absolute_percentage_error(Y_test, Y_pred)
print("MAE:", mae)
mse = mean_squared_error(Y_test, Y_pred)
print("MSE:", mse)
rmse = np.sqrt(mse)
print("RMSE:", rmse)
r2 = r2_score(Y_test, Y_pred)
print("R²:", r2)

```

MAE: 0.1682407531839539

MSE: 1623478473.0801294

RMSE: 40292.41210302666

R²: 0.2588271652481561

The statistics show that the Linear Regression Model has low performance and needs to be polished or replaced by other linear models or non-linear models to improve prediction performance.